

Compute Resources

User Agreement

<https://washu.atlassian.net/wiki/spaces/RUD/pages/1864892726/Docker+Basics+Building+Tagging+Pushing+A+Custom+Docker+Image?atlOrigin=eyJpIjoiMTVjMjNIM>

RFdiffusion (<https://github.com/RosettaCommons/RFdiffusion/tree/main>)

Mounts

Docker Image

This information should be space separated in the field that says "mounts" on the Open OnDemand Portal.

Fill out the rest of the fields with the appropriate information (explained in the quick start).

Launch the job using the submit button as described in the quick start.

Once in an interactive relion session using the following command:

We have provided the model files at `/scratch1/fs1/ris/references/RFdiffusion/models/` so the users need not download these again. These models were downloaded using the [script](#) provided in the RFdiffusion source code that shows what files to download. As of 03/13/2024 the following model files (with hashes) are downloaded and provided:

To use the provided models please add the above path to `inference.model_directory_path` configuration.

RFdiffusion may not run on thpc unless provided with a schedule path with the inference.schedule_directory_path config option. Please note: The mentioned path does not appear to need anything in it to run, but RIS is not support for RFdiffusion and cannot speculate on the config setting's use beyond its need to be explicitly defined in this context.

Batch Job Examples

Refer below is an example bsub job:

```
\ -q general \ -a 'docker(ghcr.io/washu-it-ris/rfdiffusion:latest)' \ python3.9 /app/Rfdiffusion/scripts/run_inference.py \
inference.output_prefix=$HOME/outputs/motifscaffolding \ inference.model_directory_path=/scratch1/fs1/ris/references/Rfdiffusion/models \
inference.input_pdb=/scratch1/fs1/ris/references/Rfdiffusion/inputs/5TPN.pdb \ inference.num_designs=3 \
inference.schedule_directory_path=$HOME/schedules \ contigmap.contigs='[10-40/A163-181/10-40]' \]]>
```

To use GPUs change the bsub arguments:

```
\ -q general \ -gpu 'num=1' \ -oo path/to/output \ -ls -a 'docker(ghcr.io/washu-it-ris/rfdiffusion:test)' \ python3.9 /app/Rfdiffusion/scripts/run_inference.py \
inference.output_prefix=$HOME/outputs/motifscaffolding \ inference.model_directory_path=/scratch1/fs1/ris/references/Rfdiffusion/models \
inference.input_pdb=/scratch1/fs1/ris/references/Rfdiffusion/inputs/5TPN.pdb \ inference.num_designs=3 \
inference.schedule_directory_path=$HOME/schedules \ contigmap.contigs='[10-40/A163-181/10-40]']]]>
```

Please do not forget to change <your_group> with your compute allocation group.

Update the cpumicro architecture 'some_arch' line from: cascadelake, skylake, zen2, etc.

Available Versions

Current Version:

ghcr.io/washu-it-ris/rfdiffusion

latest