11falsenonelisttrue

Compute Resources

Have questions or need help with compute, including activation or issues? Follow this link.

User Agreement

Docker Usage

The information on this page assumes that you have a knowledge base of using Docker to create images and push them to a repository for use. If you need to review that information, please see the links below.

https://washu.atlassian.net/wiki/spaces/RUD/pages/1705115761/Docker+and+the+RIS+Compute1+Platform?atlOrigin=eyJpIjoiNzc4YTZjNjIxYmQwNGI3OTk4M2Q0MW

https://washu.atlassian.net/wiki/spaces/RUD/pages/1864892726/Docker+Basics+Building+Tagging+Pushing+A+Custom+Docker+Image?atlOrigin=eyJpIjoiMTVjMjNlM

Overview

This documentation will guide you on making sure you're using the most appropriate CUDA version for your Docker image in regards to the Scientific Compute Platform.

Using the Correct Version

Examples of appropriate base images:

Nvidia has a lot of base images to develop from and can be found here: https://hub.docker.com/r/nvidia/cuda/tags

Testing Your Image

Shown below are the steps to run a test job.

Start up an interactive job with your Docker image.

There is a test script in https://github.com/WashU-IT-RIS/docker-osu-micro-benchmarks.git for the OSU GPU test.

Clone the repository.

Change directory to docker-osu-micro-benchmarks.

Run an OSU Benchmark GPU test.

Replace <test> with an OSU test that you want to run. For example, osu_bw for OSU bandwidth test.

Replace <compute-group> with the compute group you are a member of.

-G]]>